Sorted Range Reporting[*]

Yakov Nekrich[†]   Gonzalo Navarro [‡]

Abstract

We consider a variant of the orthogonal range reporting problem when all points should be reported in the sorted order of their x-coordinates. We show that reporting two-dimensional points with this additional condition can be organized (almost) as efficiently as the standard range reporting. Moreover, our results generalize and improve the previously known results for the orthogonal range successor problem and can be used to obtain better solutions for some stringology problems.

1 Introduction

An orthogonal range reporting query Q on a set of d-dimensional points S asks for all points $p \in S$ that belong to the query rectangle $Q = \prod_{i=1}^{d}[a_i, b_i]$. The orthogonal range reporting problem, that is, the problem of constructing a data structure that supports such queries, was studied extensively; see for example [1]. In this paper we consider a variant of the two-dimensional range reporting in which reported points must be sorted by one of their coordinates. Moreover, our data structures can also work in the online modus: the query answering procedure reports all points from $S \cap Q$ in increasing x-coordinate order until the procedure is terminated or all points in $S \cap Q$ are output.[1]

Some simple database queries can be represented as orthogonal range reporting queries. For instance, identifying all company employees who are between 20 and 40 years old and whose salary is in the range $[r_1, r_2]$ is equivalent to answering a range reporting query $Q = [r_1, r_2] \times [20, 40]$ on a set of points with coordinates (salary, age). Then reporting employees with the salary-age range Q sorted by their salary is equivalent to a sorted range reporting query.

Furthermore, the sorted reporting problem is a generalization of the orthogonal range successor problem (also known as the range next-value problem) [15, 8, 14, 7, 21]. The answer to an orthogonal range successor query $Q = [a, +\infty] \times [c, d]$ is the point with smallest x-coordinate[2] among all points that are in the rectangle Q. The best previously known $O(n)$ space data structure for the range successor queries uses $O(n)$ space and supports queries in $O(\log n / \log \log n)$ time [21]. The fastest previously described structure supports range successor queries in $O(\log \log n)$ time but needs $O(n \log n)$ space. In this paper we show that these results can be significantly improved.

In Section 3 we describe two data structures for range successor queries. The first structure needs $O(n)$ space and answers queries in $O(\log^{\varepsilon} n)$ time; henceforth ε denotes an arbitrarily small positive

[1]We can get increasing/decreasing x/y-coordinate ordering via coordinate changes.

[2]Previous works (e.g., [8, 21]) use slightly different definitions, but all of them are equivalent up to a simple change of coordinate system or reduction to rank space [11].

constant. The second structure needs $O(n \log \log n)$ space and supports queries in $O((\log \log n)^2)$ time. Both data structures can be used to answer sorted reporting queries in $O((k+1) \log^\varepsilon n)$ and $O((k+1)(\log \log n)^2)$ time, respectively, where k is the number of reported points. In Sections 4 and 5 we further improve the query time and describe a data structure that uses $O(n \log^\varepsilon n)$ space and supports sorted reporting queries in $O(\log \log n + k)$ time. As follows from the reduction of [17] and the lower bound of [19], any data structure that uses $O(n \log^{O(1)} n)$ space needs $\Omega(\log \log n + k)$ time to answer (unsorted) orthogonal range reporting queries. Thus we achieve optimal query time for the sorted range reporting problem. We observe that the currently best data structure for unsorted range reporting in optimal time [5] also uses $O(n \log^\varepsilon n)$ space. In Section 6 we discuss applications of sorted reporting queries to some problems related to text indexing and some geometric problems.

Our results are valid in the word RAM model. Unless specified otherwise, we measure the space usage in words of $\log n$ bits. We denote by $p.x$ and $p.y$ the coordinates of a point p. We assume that points lie on an $n \times n$ grid, i.e., that point coordinates are integers integers in $[1, n]$. We can reduce the more general case to this one by reduction to rank space [11]. The space usage will not change and the query time would increase by an additive factor $pred(n)$, where $pred(n)$ is the time needed to search in a one-dimensional set of integers [20, 19].

2   Compact Range Trees

The range tree is a handbook data structure frequently used for various orthogonal range reporting problems. Its leaves contain the x-coordinates of points; a set $S(v)$ associated with each node v contains all points whose x-coordinates are stored in the subtree rooted at v. We will assume that points of $S(v)$ are sorted by their y-coordinates. $S(v)[i]$ will denote the i-th point in $S(v)$; $S(v)[i..j]$ will denote the sorted list of points $S(v)[i], S(v)[i+1], \ldots, S(v)[j]$.

A standard range tree uses $O(n \log n)$ space, but this can be reduced by storing compact representations of sets $S(v)$. We will need to support the following two operations on compact range trees. Given a range $[c, d]$ and a node v, $noderange(c, d, v)$ finds the range $[c_v, d_v]$ such that $p.y \in [c, d]$ if and only if $p \in S(v)[c_v..d_v]$ for any $p \in S(v)$. Given an index i and a node v, $point(v, i)$ returns the coordinates of point $S(v)[i]$.

Lemma 1 *[6, 5] There exists a compact range tree that uses $O(nf(n))$ space and supports operations $point(v, i)$ and $noderange(c, d, v)$ in $O(g(n))$ and $O(g(n) + \log \log n)$ time, respectively, for (i) $f(n) = O(1)$ and $g(n) = O(\log^\varepsilon n)$; (ii) $f(n) = O(\log \log n)$ and $g(n) = O(\log \log n)$; (iii) $f(n) = O(\log^\varepsilon n)$ and $g(n) = O(1)$.*

Proof: We can support $point(v, i)$ in $O(g(n))$ time using $O(nf(n))$ space as in variants (i) and (iii) using a result from Chazelle [6]; we can support $point(v, i)$ in $O(\log \log n)$ time and $O(n \log \log n)$ space using a result from Chan et al. [5]. In the same paper [5, Lemma 2.4], the authors also showed how to support $noderange(c, d, i)$ in $O(g(n) + \log \log n)$ time and $O(n)$ additional space using a data structure that supports $point(v, i)$ in $O(g(n))$ time. $\square$

3   Sorted Reporting in Linear Space

In this section we show how a range successor query $Q = [a, +\infty] \times [c, d]$ can be answered efficiently. We combine the recursive approach of the van Emde Boas structure [20] with compact structures

for range maxima queries. A combination of succinct range minima structures and range trees was also used in [5]. A novel idea that distinguishes our data structure from the range reporting structure in [5], as well as from the previous range successor structures, is binary search on tree levels originally designed for one-dimensional searching [20]. We essentially perform a one-dimensional search for the successor of a and answer range maxima queries at each step. Let T_x denote the compact range tree on the x-coordinates of points. T_x is implemented as in variant (i) of Lemma 1; hence, we can find the interval $[c_v, d_v]$ for any node v in $O(\log^\varepsilon n)$ time. We also store a compact structure for range maximum queries $M(v)$ in every node v: given a range $[i, j]$, $M(v)$ returns the index $i \leq t \leq j$ of the point p with the greatest x-coordinate in $S(v)[i..j]$. We also store a structure for range minimum queries $M'(v)$. $M(v)$ and $M'(v)$ use $O(n)$ bits and answer queries in $O(1)$ time [9]. Hence all $M(u)$ and $M'(u)$ for $u \in T_x$ use $O(n)$ space. Finally, an $O(n)$ space level ancestor structure enables us to find the depth-d ancestor of any node $u \in T_x$ in $O(1)$ time [2].

Let π denote the search path for a in the tree T_x: π connects the root of T_x with the leaf that contains the smallest value $a_x \geq a$. Our procedure looks for the lowest node v_f on π such that $S(v) \cap Q \neq \emptyset$. For simplicity we assume that the length of π is a power of 2. We initialize v_l to the leaf that contains a_x; we initialize v_u to the root node. The node v_f is found by a binary search on π. We say that a node w is the middle node between u and v if w is on the path from u to v and the length of the path from u to w equals to the length of the path from w to v. We set the node v_m to be the middle node between v_u and v_l. Then we find the index t_m of the maximal element in $S(v_m)[c_{v_m}..d_{v_m}]$ and the point $p_m = S(v_m)[t_m]$. If $p_m.x \geq a$, then v_f is either v_m or its descendant; hence, we set $v_u = v_m$. If $p_m.x < a$, then v_f is an ancestor of v_m; hence, we set $v_l = v_m$. The search procedure continues until v_u is the parent of v_m. Finally, we test nodes v_u and v_l and identify v_f (if such v_f exists).

Fact 1 *If the child v' of v_f belongs to π, then v' is the left child of v_f.*

Proof: Suppose that v' is the right child of v_f and let v'' be the sibling of v'. By definition of v_f, $Q \cap S(v') = \emptyset$. Since v' belongs to π and v'' is the left child, $p.x < a$ for all points $p \in S(v'')$. Since $S(v_f) = S(v') \cup S(v'')$, $Q \cap S(v_f) = \emptyset$ and we obtain a contradiction. $\square$

Since $v' \in \pi$ is the left child of v_f, $p.x \geq a$ for all $p \in S(v'')$ for the sibling v'' of v. Moreover, $p.x < a$ for all points $p \in S(v')[c_{v'}, d_{v'}]$ by definition of v_f. Therefore the range successor is the point with minimal x-coordinate in $S(v'')[c_{v''}..d_{v''}]$.

The search procedure visits $O(\log \log n)$ nodes and spends $O(\log^\varepsilon n)$ time in each node, thus the total query time is $O(\log^\varepsilon n \log \log n)$. By replacing $\varepsilon' < \varepsilon$ in the above construction, we obtain the following result.

Lemma 2 *There exists a data structure that uses $O(n)$ space and answers orthogonal range successor queries in $O(\log^\varepsilon n)$ time.*

If we use the compact tree that needs $\Theta(n \log \log n)$ space, then $g(n) = \log \log n$. Using the same structure as in the proof of Lemma 2, we obtain the following.

Lemma 3 *There exists a data structure that uses $O(n \log \log n)$ space and answers orthogonal range successor queries in $O((\log \log n)^2)$ time.*

Sorted Reporting Queries. We can answer sorted reporting queries by answering a sequence of range successor queries. Consider a query $Q = [a, b] \times [c, d]$. Let p_1 be the answer to the range successor query $Q_1 = [a, +\infty] \times [c, d]$. For $i \geq 2$, let p_i be the answer to the query $Q_i = [p_{i-1}.x, +\infty] \times [c, d]$. The sequence of points $p_1, \ldots p_k$ is the sequence of k leftmost points in $[a, b] \times [c, d]$ sorted by their x-coordinates. We observe that our procedure also works in the *online modus* when k is not known in advance. That is, we can output the points of $Q \cap S$ in the left-to-right order until the procedure is stopped by the user or all points in $Q \cap S$ are reported.

Theorem 1 *There exist a data structures that uses $O(n)$ space and answer sorted range reporting queries in $O((k+1) \log^\varepsilon n)$ time, and that use $O(n \log \log n)$ space and answer those queries in $O((k+1)(\log \log n)^2)$ time.*

4 Three-Sided Reporting in Optimal Time

In this section we present optimal time data structures for two special cases of sorted two-dimensional queries. In the first part of this section we describe a data structure that answers sorted one-sided queries: for a query c we report all points p, $p.y \leq c$, sorted in increasing order of their x-coordinates. Then we will show how to answer three-sided queries, i.e., to report all points p, $a \leq p.x \leq b$ and $p.y \leq c$, sorted in increasing order by their x-coordinates.

One-Sided Sorted Reporting. We start by describing a data structure that answers queries in $O(\log n + k)$ time; our solution is based on a standard range tree decomposition of the query interval $[1, c]$ into $O(\log n)$ intervals. Then we show how to reduce the query time to $O(k + \log \log n)$. This improvement uses an additional data structure for the case when $k \leq \log n$ points must be reported.

We construct a range tree on the y-coordinates. For every node $v \in T$, the list $L(v)$ contains all points that belong to v sorted by their x-coordinates. Suppose that we want to return k points p with smallest x-coordinates such that $p.y \leq c$. We can represent the interval $[1, c]$ as a union of $O(\log n)$ node ranges for nodes $v_i \in T$. The search procedure visits each v_i and finds the leftmost point (that is, the first point) in every list $L(v_i)$. Those points are kept in a data structure D. Then we repeat the following step k times: We find the leftmost point p stored in D, output p and remove it from D. If p belongs to a list $L(v_i)$, we find the point p' that follows p in $L(v_i)$ and insert p' into D. As D contains $O(\log n)$ points, we support updates and find the leftmost point in D in $O(1)$ time [10]. Hence, we can initialize D in $O(\log n)$ time and then report k points in $O(k)$ time.

We can reduce the query time to $O(k + \log \log n)$ by constructing additional data structures. If $k \geq \log n$ the data structure described above already answers a query in $O(k + \log n) = O(k)$ time. The case $k \leq \log n$ can be handled as follows. We store for each $p \in S$ a list $V(p)$. Among all points $p' \in S$ such that $p'.y \leq p.y$ the list $V(p)$ contains $\log n$ points with the smallest x-coordinates. Points in $V(p)$ are sorted in increasing order by their x-coordinates. To find k leftmost points in $[1, c]$ for $k < \log n$, we identify the highest point $p_c \in S$ such that $p_c.y \leq c$ and report the first k points in $V(p_c)$. The point p_c can be found in $O(\log \log n)$ time using the van Emde Boas data structure [20]. If p_c is known, then a query can be answered in $O(k)$ time for any value of k.

One last improvement will be important for the data structure of Lemma 5. Let S_m denote the set of $\lceil \log \log n \rceil$ lowest points in S. We store the y-coordinates of $p \in S_m$ in the q-heap F. Using F, we can find the highest $p_m \in S_m$, such that $p_m.y \leq c$, in $O(1)$ time [10]. Let $n_c = |\{p \in S \mid p.y \leq c\}|$. If $n_c \leq \log \log n$, then $p_m = p_c$. As described above, we can answer a

query in $O(k)$ time when p_c is known. Hence, a query can be answered in $O(k)$ time if $n_c \leq \log \log n$.

Lemma 4 *There exists an $O(n \log n)$ space data structure that supports one-sided sorted range reporting queries in $O(\log \log n + k)$ time. If the highest point p with $p.y \leq c$ is known, then one-sided sorted queries can be supported in $O(k)$ time. If $|\{ p \in S \,|\, p.y \leq c \}| \leq \log \log n$, a sorted range reporting query $[1, c]$ can be answered in $O(k)$ time.*

Three-Sided Sorted Queries. We construct a range tree on x-coordinates of points. For any node v, the data structure $D(v)$ of Lemma 4 supports one-sided queries on $S(v)$ as described above. For each root-to-leaf path π we store two data structures, $R_1(\pi)$ and $R_2(\pi)$. Let π^+ and π^- be defined as follows. If v belongs to a path π and v is the left child of its parent, then its sibling v' belongs to π^+. If v belongs to π and v is the right child of its parent, then its sibling v' belongs to π^-. The data structure $R_1(\pi)$ contains the lowest point in $S(v')$ for each $v' \in \pi^+$; if $v \in \pi$ is a leaf, $R_1(\pi)$ also contains the point stored in v. The data structure $R_2(\pi)$ contains the lowest point in $S(v')$ for each $v' \in \pi^-$; if $v \in \pi$ is a leaf, $R_2(\pi)$ also contains the point stored in v. Let $lev(v)$ denote the level of a node v (the level of a node v is the length of the path from the root to v). If a point $p \in R_i(\pi)$, $i = 1, 2$, comes from a node v, then $lev(p) = lev(v)$. For a given query (c, l) the data structure $R_1(\pi)$ $(R_2(\pi))$ reports points p such that $p.y \leq c$ and $lev(p) \geq l$ sorted in decreasing (increasing) order by $lev(p)$. Since a data structure $R_i(\pi)$, $i = 1, 2$, contain $O(\log n)$ points, the point with the k-th largest (smallest) value of $lev(p)$ among all p with $p.y \leq c$ can be found in $O(1)$ time. The implementation of structures $R_i(\pi)$ is based on standard bit techniques and will be described in the full version.

Consider a query $Q = [a, b] \times [1, c]$. Let π_a and π_b be the paths from the root to a and b respectively. Suppose that the lowest node $v \in \pi_a \cap \pi_b$ is situated on level $lev(v) = l$. Then all points p such that $p.x \in [a, b]$ belong to some node v such that $v \in \pi_a^+$ and $lev(v) > l$ or $v \in \pi_b^-$ and $lev(v) > l$. We start by finding the leftmost point p in $R_1(\pi_a)$ such that $lev(p) > l$ and $p.y \leq c$. Since the x-coordinates of points in $R_1(\pi_a)$ decrease as $lev(p)$ increases, this is equivalent to finding the point $p_1 \in R_1(\pi_a)$ such that $p_1.y \leq c$ and $lev(p_1)$ is maximal. If $lev(p_1) > l$, we visit the node $v_1 \in \pi_a^+$ that contains p_1; using $D(v_1)$, we report the k leftmost points $p' \in S(v_1)$ such that $p'.y \leq c$. Then, we find the point p_2 with the next largest value of $lev(p)$ among all $p \in R_1(\pi_a)$ such that $p.y \leq c$; we visit the node $v_2 \in \pi_a^+$ that contains p_2 and proceed as above. The procedure continues until k points are output or there are no more points $p \in R_1(\pi_a)$, $lev(p) > l$ and $p.y \leq c$. If $k' < k$ points were reported, we visit selected nodes $u \in \pi_b^-$ and report remaining $k - k'$ points using a symmetric procedure.

Let k_i denote the number of reported points from the set $S(v_i)$ and let $m_i = Q \cap S(v_i)$. We spend $O(k_i)$ time in a visited node v_i if $k_i \geq \log \log n$ or $m_i < \log \log n$. If $k_j < \log \log n$ and $m_j \geq \log \log n$, then we spend $O(\log \log n + k_j)$ time in the respective node v_j. Thus we spend $O(\log \log n + k_j)$ time in a node v_j only if $m_j > k_j$, i.e., only if not all points from $S(v_j) \cap Q$ are reported. Since at most one such node v_j is visited, the total time needed to answer all one-sided queries is $O(\sum_i k_i + \log \log n) = O(\log \log n + k)$.

Lemma 5 *There exists an $O(n \log^2 n)$ space data structure that answers three-sided sorted reporting queries in $O(\log \log n + k)$ time.*

Online queries. We assumed in Lemmas 4 and 5 that parameter k is fixed and given with the query. Our data structures can also support queries in the online modus using the method originally

5

described in [3]. The main idea is that we find roughly $\Theta(k_i)$ leftmost points from the query range for $k_i = 2^i$ and $i = 1, 2, \ldots$; while k_i points are reported, we simultaneously compute the following $\Theta(k_{i+1})$ points in the background. For a more extensive description, refer to [18, Section 4.1], where the same method for a slightly different problem is described.

5 Two-Dimensional Range Reporting in Optimal Time

We store points in a compact range tree T_y on y-coordinates. We use the variant (iii) of Lemma 1 that uses $O(n \log^\varepsilon n)$ space and retrieves the coordinates of the r-th point from $S(v)$ in $O(1)$ time. Moreover, the sets $S(v)$, $v \in T_y$, are divided into groups $G_i(v)$. Each $G_i(v)$, except of the last one, contains $\lceil \log^3 n \rceil$ points. For $i < j$, each point assigned to $G_i(v)$ has smaller x-coordinate than any point in $G_j(v)$. The set $S'(v)$ contains selected elements from $S(v)$. If v is the right child of its parent, then $S'(v)$ contains $\lceil \log \log n \rceil$ points with smallest y-coordinates from each group $G_i(v)$; structure $D'(v)$ supports three-sided sorted queries of the form $[a, b] \times [0, c]$ on points of $S'(v)$. If v is the left child of its parent, then $S'(v)$ contains $\lceil \log \log n \rceil$ points with largest y-coordinates from each group $G_i(v)$; data structure $D'(v)$ supports three-sided sorted queries of the form $[a, b] \times [c, +\infty]$ on points of $S'(v)$. For each point $p' \in S'(v)$ we store the index i of the group $G_i(v)$ that contains p. We also store the point with the largest x-coordinate from each $G_i(v)$ in a structure $E(v)$ that supports $O(\log \log n)$ time searches [20].

For all points in each group $G_i(v)$ we store an array $A_i(v)$ that contains points sorted by their y-coordinates. Each point is specified by the rank of its x-coordinate in $G_i(v)$; so each entry uses $O(\log \log n)$ bits of space.

To answer a query $Q = [a, b] \times [c, d]$, we find the lowest common ancestor v_c of the leaves that contain c and d. Let v_l and v_r be the left and the right children of v_c. All points in $Q \cap S$ belong to either $([a, b] \times [c, +\infty]) \cap S(v_l)$ or $([a, b] \times [0, d]) \cap S(v_r)$. We generate the sorted list of k leftmost points in $Q \cap S$ by merging the lists of k leftmost points in $([a, b] \times [c, +\infty]) \cap S(v_l)$ and $([a, b] \times [0, d]) \cap S(v_r)$. Thus it suffices to answer sorted three-sided queries $([a, b] \times [c, +\infty])$ and $([a, b] \times [0, d])$ in nodes v_l and v_r respectively.

We consider a query $([a, b] \times [0, d]) \cap S(v_r)$; query $[a, b] \times [c, +\infty]$ is answered symmetrically. Assume $[a, b]$ fits into one group $G_i(v_r)$, i.e., all points p such that $a \le p.x \le b$ belong to one group $G_i(v_r)$. We can find the y-rank d_r of the highest point $p \in G_i(v_r)$, such that $p.y \le d$ in $O(\lg \lg n)$ time by binary search in $A_i(v_r)$. Let a_r and b_r be the ranks of a and b in $G_i(v_r)$. We can find the positions of k leftmost points in $([a_r, b_r] \times [0, d_r]) \cap G_i(v_r)$ using a data structure $H_i(v_r)$. $H_i(v_r)$ contains the y-ranks and x-ranks of points in $G_i(v_r)$ and answers sorted three-sided queries on $G_i(v_r)$. By Lemma 5, $H_i(v_r)$ uses $O(|G_i(v_r)|(\log \log n)^3)$ bits and supports queries in $O(\log \log \log n + k)$ time. Actual coordinates of points can be obtained from their ranks in $G_i(v_r)$ in $O(1)$ time per point: if the x-rank of a point is known, we can compute its position in $S(v_r)$; we obtain x-coordinates of the i-th point in $S(v_r)$ using variant (iii) of Lemma 1.

Now assume $[a, b]$ spans several groups $G_i(v_r), \ldots, G_j(v_r)$ for $i < j$. That is, the x-coordinates of all points in groups $G_{i+1}(v_r), \ldots, G_{j-1}(v_r)$ belong to $[a, b]$; the x-coordinate of at least one point in $G_i(v_r)$ ($G_j(v_r)$) is smaller than a (larger than b) but the x-coordinate of at least one point in $G_i(v_r)$ and $G_j(v_r)$ belongs to $[a, b]$. Indices i and j are found in $O(\log \log n)$ time using $E(v_r)$. We report at most k leftmost points in $([a, b] \times [0, d]) \cap G_i(v_r)$ just as described above.

Let $k_1 = |([a, b] \times [0, d]) \cap G_i(v_r)|$; if $k_1 \ge k$, the query is answered. Otherwise, we report $k' = k - k_1$ leftmost points in $([a, b] \times [0, d]) \cap (G_{i+1}(v_r) \cup \ldots \cup G_{j-1}(v_r))$ using the following

method. Let a' and b' be the minimal and the maximal x-coordinates of points in $G_{i+1}(v_r)$ and $G_{j-1}(v_r)$, respectively. The main idea is to answer the query $Q' = ([a', b'] \times [0, d]) \cap S'(v_r)$ in the online modus using the data structure $D'(v_r)$. If some group $G_t(v_r)$, $i < t < j$, contains less than $\lceil \log \log n \rceil$ points p with $p.y \leq d$, then all such p belong to $S'(v_r)$ and will be reported by $D'(v_r)$. Suppose that $D'(v_r)$ reported $\log \log n$ points that belong to the same group $G_t(v_r)$. Then we find the rank d_t of d among the y-coordinates of points in $G_t(v_r)$. Using $H_t(v_r)$, we report the positions of all points $p \in G_t(v_r)$, such that the rank of $p.y$ in $G_t(v_r)$ is at most d_t, in the left-to right order; we can also identify the coordinates of every such p in $O(1)$ time per point. The query to $H_t(v_r)$ is terminated when all such points are reported or when the total number of reported points is k.

We need $O(\log \log n + k_t)$ time to answer a query on $H_t(v_r)$, where k_t denotes the number of reported points from $G_t(v_r)$. Let $m_t = |Q' \cap G_t(v_r)|$ If G_t is the last examined group, then $k_t \leq m_t$; otherwise $k_t = m_t$. We send a query to $G_t(v_r)$ only if $G_t(v_r)$ contains at least $\log \log n$ points from Q'. Hence, a query to $G_t(v_r)$ takes $O(\log \log n + k_t) = O(k_t)$ time, unless $G_t(v_r)$ is the last examined group. Thus all necessary queries to $G_t(v_r)$ for $i<t<j$ take $O(\log \log n + k)$ time.

Finally, if the total number of points in $([a, b] \times [0, d]) \cap (G_i(v_r) \cup \ldots \cup G_{j-1}(v_r))$ is smaller than k, we also report the remaining points from $([a, b] \times [0, d]) \cap G_j(v_r)$.

The compact tree T_y uses $O(n \log^\varepsilon n)$ words of space. A data structure $D'(v)$ uses $O(|S'(v)| \log^2 n \log \log n) = O(|S(v)| \log \log n / \log n)$ words of space. Since all sets $S(v)$, $v \in T_y$, contain $O(n \log n)$ points, all $D'(v)$ use $O(n \log \log n)$ words of space. A data structure for a group $G_i(v)$ uses $O(|G_i(v)|(\log \log n)^3)$ bits. Since all $G_i(v)$ for all $v \in T_y$ contain $O(n \log n)$ elements, data structures for all groups $G_i(v)$ use $O(n(\log \log n)^3)$ words of $\log n$ bits.

Theorem 2 *There exists a $O(n \log^\varepsilon n)$ space data structure that answers two-dimensional sorted reporting queries in $O(\log \log n + k)$ time.*

6 Applications

In this section we will describe data structures for several indexing and computational geometry problems. A text (string) T of length n is pre-processed and stored in a data structure so that certain queries concerning some substrings of T can be answered efficiently.

Preliminaries. In a suffix tree $\mathcal{T}$ for a text T, every leaf of $\mathcal{T}$ is associated with a suffix of T. If the leaves of $\mathcal{T}$ are listed from left to right, then the corresponding suffixes of T are lexicographically sorted. For any pattern P, we can find in $O(|P|)$ time the special node $v \in \mathcal{T}$, called the *locus* of P. The starting position of every suffix in the subtree of $v = \mathsf{locus}(P)$ is the location of an occurrence of P. We define the rank of a suffix Suf as the number of T's suffixes that are lexicographically smaller than or equal to Suf. The ranks of all suffixes in $v = \mathsf{locus}(P)$ belong to an interval $[left(P), right(P)]$, where $left(P)$ and $right(P)$ denote the ranks of the leftmost and the rightmost suffixes in the subtree of v. Thus for any pattern P there is a unique range $[left(P), right(P)]$; pattern P occurs at position i in T if and only if the rank of suffix $T[i..n]$ belongs to $[left(P), right(P)]$. Refer to [13] for a more extensive description of suffix trees and related concepts.

We will frequently use a special set of points, further called *the position set for T*. Every point p in the position set corresponds to a unique suffix Suf of a string T; the y-coordinate of p equals to the rank of Suf and the x-coordinate of p equals to the starting position of Suf in T.

Successive List Indexing. In this problem a query consists of a pattern P and an index j, $1 \leq j \leq n$. We want to find the first (leftmost) occurrence of P at position $i \geq j$. A successive list indexing query (P, j) is equivalent to finding the point p from the position set such that p belongs to the range $[j, n] \times [left(P), right(P)]$ and the x-coordinate of p is minimal. Thus a list indexing query is equivalent to a range successor query on the position set. Using Theorems 1 and 2 to answer range successor queries, we obtain the following result.

Corollary 1 *We can store a string T in an $O(nf(n))$ space data structure, so that for any pattern P and any index j, $1 \leq j \leq n$, the leftmost occurrence of P at position $i \geq j$ can be found in $O(g(n))$ time for (i) $f(n) = O(1)$ and $g(n) = O(\log^\varepsilon n)$; (ii) $f(n) = O(\log\log n)$ and $g(n) = O((\log\log n)^2)$; (iii) $f(n) = O(\log^\varepsilon n)$ and $g(n) = O(\log\log n)$.*

Range Non-Overlapping Indexing. In the string statistics problem we want to find the maximum number of non-overlapping occurrences of a pattern P. In [14] the *range non-overlapping indexing problem* was introduced: instead of just computing the maximum number of occurrences we want to find the longest sequence of non-overlapping occurrences of P. It was shown [14] that the range non-overlapping indexing problem can be solved via k successive list indexing queries; here k denotes the maximal number of non-overlapping occurrences.

Corollary 2 *The range non-overlapping indexing problem can be solved in $O(|P| + kg(n))$ time with an $O(nf(n))$ space data structure, where $g(n)$ and $f(n)$ are defined as in Corollary 1.*

Other, more far-fetched applications, are described next.

6.1   Pattern Matching with Variable-Length Don't Cares

We must determine whether a query pattern $P = P_1 * P_2 * P_3 \ldots * P_m$ occurs in T. The special symbol $*$ is the Kleene star symbol; it corresponds to an arbitrary sequence of (zero or more) characters from the original alphabet of T. The parameter m can be specified at query time. In [22] the authors showed how to answer such queries in $O(\sum_{i=1}^m |P_i|)$ and $O(n)$ space in the case when the alphabet size is $\log^{O(1)} n$. In this paper we describe a data structure for an arbitrarily large alphabet. Using the approach of [22], we can reduce such a query for P to answering m successive list indexing queries. First, we identify the leftmost occurrence of P_1 in T by answering the successive list indexing query $(P_1, 1)$. Let j_1 denote the leftmost position of P_1. $P_1 * P_2 * P_3 \ldots * P_m$ occurs in T if and only if $P_2 * P_3 \ldots * P_m$ occurs at position $i \geq j_1 + |P_1|$. We find the leftmost occurrence $j_2 \geq j_1 + |P_1|$ of P_2 by answering the query $(P_2, j_1 + |P_1|)$. $P_2 * P_3 \ldots * P_m$ occurs in T at position $i_2 \geq j_1 + |P_1|$ if and only if $P_3 * P_m$ occurs at position $i_3 \geq j_2 + |P_2|$. Proceeding in the same way we find the leftmost possible positions for $P_4 * \ldots * P_m$. Thus we answer m successive list indexing queries (P_t, i_t), $t = 1, \ldots, m$; here $i_1 = 1$, $i_t = j_{t-1} + |P_{i-1}|$ for $t \geq 2$, and j_{t-1} denotes the answer to the $(t-1)$-th query.

Corollary 3 *We can determine whether a text T contains a substring $P = P_1 * \ldots P_{m-1} * P_m$ in $O(\sum_{i=1}^m |P_i| + mg(n))$ time using an $O(nf(n))$ space data structure, where $g(n)$ and $f(n)$ are defined as in Corollaries 1 and 2.*

6.2 Ordered Substring Searching

Suppose that a data structure contains a text T and we want to report occurrences of a query pattern P in the left-to-right order, i.e., in the same order as they appear in T; in some case we may want to find only the k leftmost occurrences. In this section we describe two solutions for this problem. Then we show how sorted range reporting can be used to solve the position-restricted variant of this problem. We denote by occ the number of P's occurrences in T that are reported when a query is answered.

Data Structure with Optimal Query Time. Such queries can be answered in $O(|P| + \text{occ})$ time and $O(n)$ space using the suffix tree and the data structure of Brodal *et al.* [3]. Positions of suffixes are stored in lexicographic order in the suffix array A; the k-th entry $A[k]$ contains the starting position of the k-th suffix in the lexicographic order. In [3] the authors described an $O(n)$ space data structure that answers online sorted range reporting queries: for any $i \geq j$, we can report in $O(j-i+1)$ time all entries $A[t]$, $i \leq t \leq j$, sorted in increasing order by their values. Occurrences of a pattern P can be reported in the left-to-right order as follows. Using a suffix tree, we find $left(P)$ and $right(P)$ in $O(|P|)$ time. Then we report all suffixes in the interval $[left(P), right(P)]$ sorted by their starting positions using the data structure of [3] on A.

Corollary 4 *We can answer a sorted substring matching query in $O(|P|+\text{occ})$ time using a $O(n)$ space data structure*

Succinct Data Structure. The space usage of a data structure for sorted pattern matching can be further reduced. We store a compressed suffix array for T and a succinct data structure for range minimum queries. We use the implementation of the compressed suffix array described in [12] that needs $(1+1/\varepsilon)nH_k + o(n)$ bits for $\sigma = \log^{O(1)} n$, where σ denotes the alphabet size and H_k is the k-th order entropy. Using the results of [12], we can find the position of the i-th lexicographically smallest suffix in $O(\log^\varepsilon n)$ time. We can also find $left(P)$ and $right(P)$ for any P in $O(|P|)$ time. We also store the range minimum data structure [9] for the array A defined above. For any $i \leq j$, we can find such $k = \mathsf{rmq}(i,j)$ that $A[k] \leq A[t]$ for any $i \leq t \leq j$. Observe that A itself is not stored; we only store the structure from [9] that uses $O(n)$ bits of space. Now occurrences of P are reported as follows. An initially empty queue Q contains suffix positions; with every suffix position p we also store an interval $[l_p, r_p]$ and the rank i_p of the suffix that starts at position p. Let $l = left(P)$ and $r = right(P)$. We find $i_f = \mathsf{rmq}(l,r)$ and the position p_f of the suffix with rank i_f. The position p_f with its rank i_f and the associated interval $[l, r]$ is inserted into Q. We repeat the following steps until Q is empty. The item with the minimal value of p_t is extracted from Q. Let i_t and $[l_t, r_t]$ denote the rank and interval stored with p_t. We answer queries $i' = \mathsf{rmq}(l_t, i_t - 1)$ and $i'' = \mathsf{rmq}(i_t + 1, r_t)$ and identify the positions p', p'' of suffixes with ranks i', i''. Finally, we insert items $(p', i', [l_t, i_t - 1])$ and $(p'', i'', [i_t + 1, r_t])$ into Q. Using the van Emde Boas data structure, we can implement each operation on Q in $O(\log \log n)$ time. We can find the position of a suffix with rank i in $O(\log^\varepsilon n)$ time. Thus the total time that we need to answer a query is $O(|P|+\text{occ} \log^\varepsilon n)$. Our data structure uses $(1+1/\varepsilon)nH_k+O(n)$ bits. We observe however that we need $O(\text{occ} \log n)$ additional bits at the time when a query is answered.

Corollary 5 *If the alphabet size $\sigma = \log^{O(1)} n$, then we can answer an ordered substring searching query in $O(|P| + \text{occ} \log^\varepsilon n)$ time using a $(1+1/\varepsilon)nH_k + O(n)$-bit data structure*

Position-Restricted Ordered Substring Searching The position restricted substring searching problem was introduced by Mäkinen and Navarro in [16]. Given a range $[i, j]$ we want to report all occurrences of P that start at position t, $i \leq t \leq j$. If we want to report occurrences of P at positions from $[i, j]$ in the sorted order, then this is equivalent to answering a sorted range reporting query $[i, j] \times [left(P), right(P)]$. Hence, we can obtain the same time-space trade-offs as in Theorems 1 and 2.

6.3 Maximal Points in a 2D Range and Rectangular Visibility

A point p *dominates* another point q if $p.x \geq q.x$ and $p.y \geq q.y$. A point $p \in S$ is *a maximal point* if p is not dominated by any other point $q \in S$. In a two-dimensional maximal points range query, we must find all maximal points in $Q \cap S$ for a query rectangle Q. We refer to [4] and references therein for description of previous results.

We can answer such queries using orthogonal range successor queries. For simplicity, we assume that all points have different x- and y-coordinates. Suppose that maximal points in the range $Q = [a, b] \times [c, d]$ must be listed. For $i \geq 1$, we report a point p_i such that $p_i.x \geq p.x$ for any $p \in Q_{i-1} \cap S$, where $Q_0 = Q$ and $Q_j = [a, p_i.x] \times [p_i.y, d]$ for $j \geq 1$. Our reporting procedure is completed when $Q_i \cap S = \emptyset$. Clearly, finding a point p_i or determining that no such p_i exists is equivalent to answering a range successor query for Q_{i-1}. Thus we can find k maximal points in $O(kg(n))$ time using an $O(nf(n))$ space data structure, where $g(n)$ and $f(n)$ are again defined as in Corollary 1.

A point $p \in S$ is rectangularly visible from a point q if $Q_{pq} \cap S = \emptyset$, where Q_{pq} is the rectangle with points p and q at its opposite corners. In the rectangle visibility problem, we must determine all points $p \in S$ that are visible from a query point q. Rectangular visibility problem is equivalent to finding maximal points in $Q \cap S$ for $Q = [0, q.x] \times [0, q.y]$. Hence, we can find points rectangularly visible from q in $O(kg(n))$ time using an $O(nf(n))$ space data structure.

References

[1] P. Agarwal and J. Erickson. Geometric range searching and its relatives. *Contemporary Mathematics*, 223:1–56, 1999.

[2] M. A. Bender and M. Farach-Colton. The level ancestor problem simplified. *Theor. Comput. Sci.*, 321(1):5–12, 2004.

[3] G. S. Brodal, R. Fagerberg, M. Greve, and A. López-Ortiz. Online sorted range reporting. In *Proc. 20th ISAAC*, pages 173–182, 2009.

[4] G. S. Brodal and K. Tsakalidis. Dynamic planar range maxima queries. In *Proc. 38th ICALP*, pages 256–267, 2011.

[5] T. M. Chan, K. G. Larsen, and M. Patrascu. Orthogonal range searching on the RAM, revisited. In *Proc. 27th SoCG*, pages 1–10, 2011.

[6] B. Chazelle. A functional approach to data structures and its use in multidimensional searching. *SIAM J. Comput.*, 17(3):427–462, 1988.

[7] M. Crochemore, C. S. Iliopoulos, M. Kubica, M. S. Rahman, and T. Walen. Improved algorithms for the range next value problem and applications. In *Proc. 25th STACS*, pages 205–216, 2008.

[8] M. Crochemore, C. S. Iliopoulos, and M. S. Rahman. Finding patterns in given intervals. In *Proc. 32nd MFCS*, pages 645–656, 2007.

[9] J. Fischer. Optimal succinctness for range minimum queries. In *Proc. 9th LATIN*, pages 158–169, 2010.

[10] M. L. Fredman and D. E. Willard. Trans-dichotomous algorithms for minimum spanning trees and shortest paths. *J. Comput. Syst. Sci.*, 48(3):533–551, 1994.

[11] H. N. Gabow, J. L. Bentley, and R. E. Tarjan. Scaling and related techniques for geometry problems. In *Proc. 16th STOC*, pages 135–143, 1984.

[12] R. Grossi, A. Gupta, and J. S. Vitter. High-order entropy-compressed text indexes. In *Proc. 14th SODA*, pages 841–850, 2003.

[13] D. Gusfield. *Algorithms on Strings, Trees, and Sequences - Computer Science and Computational Biology.* Cambridge University Press, 1997.

[14] O. Keller, T. Kopelowitz, and M. Lewenstein. Range non-overlapping indexing and successive list indexing. In *Proc. 10th WADS*, pages 625–636, 2007.

[15] H.-P. Lenhof and M. H. M. Smid. Using persistent data structures for adding range restrictions to searching problems. *RAIRO Theor. Inf. and Appl.*, 28(1):25–49, 1994.

[16] V. Mäkinen and G. Navarro. Rank and select revisited and extended. *Theor. Comput. Sci.*, 387(3):332–347, 2007.

[17] P. B. Miltersen, N. Nisan, S. Safra, and A. Wigderson. On data structures and asymmetric communication complexity. *J. Comput. Syst. Sci.*, 57(1):37–49, 1998.

[18] G. Navarro and Y. Nekrich. Top-k document retrieval in optimal time and linear space. In *Proc. 22nd SODA*, pages 1066–1077, 2012.

[19] M. Pătraşcu and M. Thorup. Time-space trade-offs for predecessor search. In *Proc. 38th STOC*, pages 232–240, 2006.

[20] P. van Emde Boas, R. Kaas, and E. Zijlstra. Design and implementation of an efficient priority queue. *Math. Sys. Theory*, 10:99–127, 1977.

[21] C.-C. Yu, W.-K. Hon, and B.-F. Wang. Improved data structures for the orthogonal range successor problem. *Comput. Geom. Theory Appl.*, 44:148–159, 2011.

[22] C.-C. Yu, B.-F. Wang, and C.-C. Kuo. Efficient indexes for the positional pattern matching problem and two related problems over small alphabets. In *Proc. 21st ISAAC*, pages 13–24, 2010.