



Two-dimensional range successor in optimal time and almost linear space

Gelin Zhou

David R. Cheriton School of Computer Science, University of Waterloo, Canada



ARTICLE INFO

Article history:

Received 30 March 2015

Accepted 5 September 2015

Available online 21 October 2015

Communicated by X. Wu

Keywords:

Data structures

Computational geometry

Range successor

Sorted range reporting

Planar orthogonal skyline reporting

ABSTRACT

In this article, we revisit the problem of supporting two-dimensional range successor queries. We present a data structure with $O(n \lg \lg n)$ words of space and $O(\lg \lg n)$ query time. This improves the work of Nekrich and Navarro (2012) by a factor of $\lg \lg n$ in query time, or a factor of $\lg^\epsilon n$ in space cost, where ϵ is an arbitrary positive constant. Our data structure matches the state of the art for two-dimensional range emptiness queries and achieves the optimal query time. This structure has fruitful applications in computational geometry and text indexing.

© 2015 Elsevier B.V. All rights reserved.

1. Introduction

The problems of supporting range aggregate queries are well-known in the communities of computational geometry and data structures. In these problems, we maintain an input point set S , so that, for any query range Q , certain functions over $S \cap Q$ can be computed efficiently. In this article we study the support for two-dimensional orthogonal *range successor* queries, or *range next-value* queries [12], for which the input point set is on the plane, and a query asks for the leftmost point in the orthogonal query range. This type of queries generalizes both *predecessor search* queries and two-dimensional orthogonal *range emptiness* queries.

As pointed out by Nekrich and Navarro [11] and Lewenstein [10], range successor queries have fruitful applications in the field of text indexing. However, only a few results have been presented for this problem [5,12,11]. The best ones have been obtained by Nekrich and Navarro [11]. Their linear space data structure requires $O(\lg^\epsilon n)$ query time, their structure using $O(n \lg \lg n)$ words of space re-

quires $O(\lg^2 \lg n)$ query time, and their structure achieving the optimal $O(\lg \lg n)$ query time occupies $O(n \lg^\epsilon n)$ words of space, where ϵ is an arbitrary positive constant.

As a variant of the classical *sorted range reporting* problem, the *sorted range reporting* problem is closely related to range successor queries. In the sorted range reporting problem, points in the orthogonal query range are reported in increasing order of their x -coordinates.¹ In addition, the query-answering procedure should work in an online fashion: points are reported in increasing order of x -coordinates until the procedure is terminated or all points in the query range are reported. This problem was proposed by Nekrich and Navarro [11]. Let k denote the number of points in the given query range. Their linear space data structure requires $O(\lg^\epsilon n + k \lg^\epsilon n)$ query time, and their data structure with the optimal $O(\lg \lg n + k)$ query time occupies $O(n \lg^\epsilon n)$ words of space. These results match the state of the art for *unsorted range reporting* [2]. However, their data structure using $O(n \lg \lg n)$ words of space requires $O(\lg^2 \lg n + k \lg^2 \lg n)$ query time,

E-mail address: g5zhou@uwaterloo.ca.

¹ Increasing/decreasing x/y -coordinate ordering can be easily supported via coordinate changes.

which is slower than the corresponding result for unsorted range reporting [2] by a factor of $\lg \lg n$.

Planar orthogonal skyline reporting queries play an important role in database systems for multi-criteria optimization [9]. As pointed out by Brodal and Larsen [1], this type of queries can be solved by calling range successor queries recursively. Thus Nekrich and Navarro's [11] range successor structures can directly support planar orthogonal skyline reporting queries within the same amount of space cost and $k + 1$ times the amount of query time, where k is the size of output. Brodal and Larsen [1] further designed two data structures based on a different idea. Their first structure occupies $O(n \lg^\epsilon n)$ words of space and answers a query within $O(\lg n / \lg \lg n + k)$ query time, and the second one supports queries using $O(\lg n / \lg \lg n + k \lg \lg n)$ query time and $O(n \lg \lg n)$ words of space.

In this article we present a word-RAM data structure for range successor queries that uses $O(n \lg \lg n)$ words of space and the optimal $O(\lg \lg n)$ query time. Previous data structures using the same amount of space require $O(\lg^2 \lg n)$ query time, and previous data structures achieving the same query time occupy $O(n \lg^\epsilon n)$ words of space [11]. Our data structure also supports sorted range reporting queries and planar orthogonal skyline reporting queries in $O(\lg \lg n + k \lg \lg n)$ time, while previous data structures using the same amount of space require $O(\lg^2 \lg n + k \lg^2 \lg n)$ query time for sorted reporting [11], and $O(\lg n / \lg \lg n + k \lg \lg n)$ or $O(\lg^2 \lg n + k \lg^2 \lg n)$ query time for skyline reporting [1,11].

The underlying computational model throughout this work is the standard word RAM model with word size $w = \Omega(\lg n)$. We assume that the given point set is in an $n \times n$ grid, or *rank space*. Every two points have different x/y -coordinates.

The rest of this article is organized as follows: In Section 2 we review Chan, Larsen and Pătraşcu's data structures [2] for unsorted range reporting and Nekrich and Navarro's results [11] for the sorted variant. In Section 3 we present our data structures for range successor queries.

2. Preliminary

For completeness, we describe Chan et al.'s work [2] for unsorted range reporting in addition to Nekrich and Navarro's work [11] for sorted range reporting. We may modify their presentations for the sake of consistency.

2.1. Unsorted range reporting

Chan et al.'s data structures [2] for range reporting are based on the *wavelet tree* [3,7]. A conceptual range tree \mathcal{T} is built on $[1..n]$, where n is assumed to be a power of two. Every node v in \mathcal{T} has an associated range. Let $S(v)$ denote the set of points whose y -coordinates are in this range; clearly $S(v)$ contains only one point for every leaf node v at the bottom of \mathcal{T} . An internal node v has two children v_ℓ and v_r , whose associated ranges form a disjoint union of v 's associated range. Points in $S(v)$ are conceptually listed as $S(v)[1], S(v)[2], \dots$ in increasing order of x -coordinates. For each of these points, we write

down a 0-bit if this point is also in $S(v_\ell)$, or a 1-bit otherwise. These bits are concatenated and maintained as a bit vector, such that rank/select operations can be performed in constant time [4].

To achieve efficient query time, Chan et al. [2] formulated the following two operations as the *ball-inheritance problem*: Given a node v in \mathcal{T} and a range $[a..b]$ on the x -axis, $\text{noderange}(v, a, b)$ returns the range $[a_v..b_v]$ such that $S(v)[a_v]$ is the first one whose x -coordinate is no smaller than a and $S(v)[b_v]$ is the last one whose x -coordinate is no greater than b . Given a node v in \mathcal{T} and an index $1 \leq i \leq |S(v)|$, $\text{point}(v, i)$ returns the coordinates of $S(v)[i]$. The following lemma addresses their results.

Lemma 2.1. (See [3,2].) *Using $O(nf(n))$ words of space, one can support $\text{noderange}(v, a, b)$ and $\text{point}(v, i)$ within $O(g(n) + \lg \lg n)$ and $O(g(n))$ query time, respectively, where*

1. $f(n) = O(1)$ and $g(n) = O(\lg^\epsilon n)$;
2. $f(n) = O(\lg \lg n)$ and $g(n) = O(\lg \lg n)$; or
3. $f(n) = O(\lg^\epsilon n)$ and $g(n) = O(1)$.

A range reporting query $Q = [a..b] \times [c..d]$ over the point set S can be answered as follows. First we find node v in \mathcal{T} , which is the lowest common ancestor of the leaf nodes that correspond to c and d . Let v_ℓ and v_r denote the children of v . It is clear that the associated range of v contains $[c..d]$, and the associated ranges of v_ℓ and v_r both intersect $[c..d]$. Therefore, $S \cap Q$ can be decomposed into $S(v_\ell) \cap ([a..b] \times [c..\infty])$ and $S(v_r) \cap ([a..b] \times [-\infty..d])$. It is sufficient to show how to support $S(v_r) \cap ([a..b] \times [-\infty..d])$ only. Setting $[a_{v_r}..b_{v_r}]$ to be $\text{noderange}(v_r, a, b)$, we only need to report the points in $S(v_r)[a_{v_r}..b_{v_r}]$ whose y -coordinates are no greater than d . To perform this step efficiently, an indexing structure for range minimum queries over $S(v_r)$ is built. Given $1 \leq i \leq j \leq |S(v)|$, this structure returns the position of the point with the smallest y -coordinate in $S(v_r)[i..j]$ using $O(1)$ time and $2|S(v_r)| + o(|S(v_r)|)$ bits of space [6]. The following procedure reports k points in $O(kg(n))$ time:

Setting $[i..j] = [a_{v_r}..b_{v_r}]$ initially, we query the smallest y -coordinate in $S(v_r)[i..j]$, which is assumed to be of index p . Using Lemma 2.1, we retrieve the y -coordinate of $S(v_r)[p]$. The procedure terminates if the y -coordinate is greater than d ; otherwise $S(v_r)[p]$ is reported and recursive calls on $[i..p-1]$ and $[p+1..j]$ are made.

The following lemma summarizes the discussions in this section.

Lemma 2.2. (See [2].) *Using $O(nf(n))$ words of space, one can support two-dimensional range reporting queries with $O(\lg \lg n + g(n) + kg(n))$ query time, where k is the output size,*

1. $f(n) = O(1)$ and $g(n) = O(\lg^\epsilon n)$;
2. $f(n) = O(\lg \lg n)$ and $g(n) = O(\lg \lg n)$; or
3. $f(n) = O(\lg^\epsilon n)$ and $g(n) = O(1)$.

Remark. It is noteworthy that the first point returned by this algorithm is the lowest point in $S(v_r) \cap Q$. However,

the first point returned by the other 3-sided query is the highest point in $S(v_\ell) \cap Q$. That being said, this algorithm cannot directly retrieve the lowest or highest point in $S \cap Q$.

2.2. Suboptimal range successor

Nekrich and Navarro [11] showed that Chan et al.'s data structures [2] could support range successor queries after certain modifications. For simplicity, x/y -coordinates are swapped and a query returns the lowest point (i.e., the one with the smallest y -coordinate) in the query range $Q = [a..b] \times [c..d]$. Let π denote the path from the root of the conceptual range tree \mathcal{T} to the leaf that corresponds to c . The basic idea is to find the lowest node v_f on π such that $S(v_f) \cap Q \neq \emptyset$. Let v denote the lowest common ancestor of the leaf nodes that correspond to c and d . It is clear that, for every node u on π that is deeper than v , its associated range contains c but not d . It implies that $S(u) \cap Q = S(u) \cap ([a..b] \times [c..\infty))$. One can determine if $S(u) \cap Q \neq \emptyset$ by determining if the 3-sided query contains any point. Therefore v_f can be computed using binary search on π . This requires to compute $O(\lg \lg n)$ 3-sided emptiness queries.

If v_f is a leaf node, then the only point in $S(v_f)$ is the answer. If v_f is an internal node, the child of v_f on π must be the left child. Otherwise the associated range of the left child would not intersect $[c..d]$, and the assumption on v_f would be contradicted. Since the left child of v_f does not correspond to any point in the query range Q , the right child must correspond to at least a point in Q . Using the algorithm described in the previous section, we can find such a point using range minimum queries. The first point returned is by chance the lowest one.

The following lemma summarizes the above discussions.

Lemma 2.3. (See [11].) *Using $O(nf(n))$ words of space, one can support two-dimensional range successor queries with $O(g(n) \lg \lg n)$ query time, where*

1. $f(n) = O(1)$ and $g(n) = O(\lg^\epsilon n)$; or
2. $f(n) = O(\lg \lg n)$ and $g(n) = O(\lg \lg n)$.

3. Range successor in optimal query time

Our data structure for range successor queries is obtained by modifying Nekrich and Navarro's [11] approach to sorted range reporting. More precisely, instead of performing binary searches on the root-to-leaf path π , we will consider and support 3-sided range successor queries, for which the leftmost point in $S \cap ([a..b] \times [-\infty..d])$ is returned. With a different partition strategy for points in $S(v)$ and appropriate use of Grossi et al.'s [8] predecessor search structure, our index for 3-sided range successor queries requires less space than that of Nekrich and Navarro. The following lemma shows a preliminary result for our data structure.

Lemma 3.1. (See Lemma 5 in [11].) *Given a set of n points, one can support 3-sided range successor queries using $O(n \lg^3 n)$ bits of space and $O(\lg \lg n)$ query time.*

Now we describe our data structures. We first build the same data structures as the second variant of Lemma 2.2. Let $Q' = [a..b] \times [-\infty..d]$ denote a 3-sided query range. We further construct auxiliary data structures on every $S(v)$ such that the leftmost point in $S(v) \cap Q'$ can be returned in $O(\lg \lg n)$ time, using $O(|S(v)| \lg \lg n)$ bits of additional space.

As we have mentioned, points in $S(v)$ are conceptually listed in increasing order of x -coordinates. These points are divided into blocks $B_1(v), B_2(v), \dots$ of size $\lceil \lg^3 n \rceil$ (the last block may contain less). $D(v)$ stores the lowest point in each block explicitly, and is maintained using Lemma 3.1 such that 3-sided range successor queries on $D(v)$ can be supported in $O(\lg \lg n)$ time. This auxiliary data structure occupies $O(|D(v)| \lg^3 |D(v)|) = O(|S(v)|)$ bits of space.

For some constant $0 < \epsilon < 1$, the points in every block $B_i(v)$ are further divided into sub-blocks $SB_{i,1}(v), SB_{i,2}(v), \dots$ of size $\lceil \lg^\epsilon n \rceil$ (the last sub-block may contain less). In addition, their x/y -coordinates are rewritten as the x/y -ranks within this block. A rank can be represented in $O(\lg \lg n)$ bits, and all the ranks require $O(|S(v)| \lg \lg n)$ bits over all blocks. $E_i(v)$ stores the lowest point in every sub-block explicitly, and is also maintained using Lemma 3.1 to support 3-sided range successor queries on $E_i(v)$ in $O(\lg \lg \lg n)$ time. This auxiliary data structure requires only $O(|E_i(v)| \lg^3 |E_i(v)|) = O(|B_i(v)| \lg^3 \lg n / \lceil \lg^\epsilon n \rceil) = o(|B_i(v)|)$ bits of additional space. Thus the space cost over all $E_i(v)$'s is $o(|S(v)|)$ bits.

Now we show how to answer the 3-sided range successor query $Q' = [a..b] \times [-\infty..d]$ over $S(v)$. First we compute $[a_v..b_v]$ using *noderange*(v, a, b), which requires $O(\lg \lg n)$ time. Let $B_{i_1}(v)$ and $B_{i_2}(v)$ be the block that contains a_v and b_v , respectively. That being said, $[a_v..b_v]$ spans over blocks $B_{i_1}(v), \dots, B_{i_2}(v)$. We only consider the case in which $i_1 < i_2$; the remaining cases can be handled similarly. We first attempt to find the leftmost point in $B_{i_1}(v) \cap Q'$ (we will show how to do it later). If such a point exists, our algorithm terminates and the point is returned. Otherwise, we query $D(v)$ to find the leftmost block among $B_{i_1+1}(v), \dots, B_{i_2-1}(v)$ that intersects Q' . Let $B_t(v)$ denote the block. We query $B_t(v) \cap Q'$ and obtain the result. If such $B_t(v)$ does not exist, we query $B_{i_2}(v) \cap Q'$.

The remaining issue is to find the leftmost point of the intersection of a single block $B_t(v)$ and the given 3-sided range Q' efficiently, for which we have the following lemma:

Lemma 3.2. *The leftmost point in $B_t(v) \cap Q'$ can be found using $O(\lg \lg n)$ time and $O(|B_t(v)| \lg \lg n)$ bits of extra space in addition to a global lookup table of size $o(n)$ bits.*

Proof. We first compute the ranks of a, b and d within this block. Let them be a', b' and d' , respectively. Here a' and b' can be computed directly from a_v and b_v in constant time. The computation of d' could be done by performing binary search on the points of $B_t(v)$. However, this

would require $O(\lg^2 \lg n)$ time, since $\text{point}(v, i)$ would be called $O(\lg \lg n)$ times. Instead, we make use of Grossi et al.'s [8, Lemma 3.3] succinct indices for the y -coordinates of points in $B_t(v)$. As mentioned in Chan et al.'s work [2], this succinct index requires $O(\lg \lg n)$ bits of extra space per point, and supports predecessor search using $O(\lg \lg n)$ time in addition to $O(1)$ calls to $\text{point}(v, i)$. Hence, the value of d' can be determined in $O(\lg \lg n)$ time.

Let $SB_{t,j_1}(v)$ and $SB_{t,j_2}(v)$ be the sub-blocks that contain a' and b' , respectively. As the procedure described previously, we only consider the case in which $j_1 < j_2$. Let Q_t denote the query $[a'..b'] \times [-\infty..d']$ within the block $B_t(v)$. We find and return the leftmost point in $SB_{t,j_1}(v) \cap Q_t$ if such a point exists. Otherwise, we query $E_t(v)$ to find the leftmost sub-block among $SB_{t,j_1+1}(v), \dots, SB_{t,j_2-1}(v)$ that intersects Q_t . Let $SB_{t,p}(v)$ denote the block; clearly the answer is the leftmost point in $SB_{t,p}(v) \cap Q_t$. If such $B_t(v)$ does not exist, we query $SB_{t,j_2}(v) \cap Q_t$. In all these cases, we need only to find the leftmost point within a sub-block. This can be done using $O(1)$ time and a global lookup table of $o(n)$ bits of additional space, since there are only $2^{\lceil \lg^\epsilon n \rceil \times O(\lg \lg n)} = O(n^{1-\delta})$ different sub-blocks, for some constant $\delta > 0$. \square

Summarizing the discussion above, we can find the leftmost point in $S(v) \cap Q'$ in $O(\lg \lg n)$ time using $O(\lg \lg n)$ bits of extra space per point and a global lookup table of size $o(n)$ bits. We also construct auxiliary data structures on every $S(v)$ for 3-sided range successor queries of the form $[a..b] \times [c..\infty]$. Combining both parts and following the approach in Section 2.1, we can find the leftmost point in a 4-sided query range in $O(\lg \lg n)$ time. Therefore we have the following theorem:

Theorem 3.3. *Under the word RAM model with word size $w = \Omega(\lg n)$, a set of n planar points in rank space can be maintained using $O(n \lg \lg n)$ words of space, so that two-dimensional range successor queries can be supported in $O(\lg \lg n)$ query time. Repeatedly using this data structure, two-dimensional sorted range reporting queries and planar orthogonal skyline reporting queries can be answered in $O(\lg \lg n + k \lg \lg n)$ time, where k is the size of output.*

References

- [1] Gerth Stølting Brodal, Kasper Green Larsen, Optimal planar orthogonal skyline counting queries, in: Algorithm Theory – SWAT 2014 – Proceedings of the 14th Scandinavian Symposium and Workshops, Copenhagen, Denmark, July 2–4, 2014, 2014, pp. 110–121.
- [2] Timothy M. Chan, Kasper Green Larsen, Mihai Pătraşcu, Orthogonal range searching on the RAM, revisited, in: Proceedings of the 27th ACM Symposium on Computational Geometry, Paris, France, June 13–15, 2011, 2011, pp. 1–10.
- [3] Bernard Chazelle, A functional approach to data structures and its use in multidimensional searching, SIAM J. Comput. 17 (3) (1988) 427–462.
- [4] David R. Clark, J. Ian Munro, Efficient suffix trees on secondary storage (extended abstract), in: Proceedings of the Seventh Annual ACM-SIAM Symposium on Discrete Algorithms, Atlanta, Georgia, 28–30 January, 1996, 1996, pp. 383–391.
- [5] Maxime Crochemore, Costas S. Iliopoulos, Marcin Kubica, Mohammad Sohel Rahman, Tomasz Walen, Improved algorithms for the range next value problem and applications, in: STACS 2008, Proceedings of the 25th Annual Symposium on Theoretical Aspects of Computer Science, Bordeaux, France, February 21–23, 2008, 2008, pp. 205–216.
- [6] Johannes Fischer, Volker Heun, Space-efficient preprocessing schemes for range minimum queries on static arrays, SIAM J. Comput. 40 (2) (2011) 465–492.
- [7] Roberto Grossi, Ankur Gupta, Jeffrey Scott Vitter, High-order entropy-compressed text indexes, in: Proceedings of the Fourteenth Annual ACM-SIAM Symposium on Discrete Algorithms, Baltimore, Maryland, USA, January 12–14, 2003, 2003, pp. 841–850.
- [8] Roberto Grossi, Alessio Orlandi, Rajeev Raman, S. Srinivasa Rao, More haste, less waste: lowering the redundancy in fully indexable dictionaries, in: Proceedings of the 26th International Symposium on Theoretical Aspects of Computer Science, STACS 2009, Freiburg, Germany, February 26–28, 2009, 2009, pp. 517–528.
- [9] Casper Kejlberg-Rasmussen, Yufei Tao, Konstantinos Tsakalidis, Kostas Tsichlas, Jeonghun Yoon, I/o-efficient planar range skyline and attrition priority queues, in: Proceedings of the 32nd ACM SIGMOD-SIGACT-SIGART Symposium on Principles of Database Systems, PODS 2013, New York, NY, USA, June 22–27, 2013, 2013, pp. 103–114.
- [10] Moshe Lewenstein, Orthogonal range searching for text indexing, in: Space-Efficient Data Structures, Streams, and Algorithms – Papers in Honor of J. Ian Munro on the Occasion of His 66th Birthday, 2013, pp. 267–302.
- [11] Yakov Nekrich, Gonzalo Navarro, Sorted range reporting, in: Algorithm Theory – SWAT 2012 – Proceedings of the 13th Scandinavian Symposium and Workshops, Helsinki, Finland, July 4–6, 2012, 2012, pp. 271–282.
- [12] Chih-Chiang Yu, Wing-Kai Hon, Bing-Feng Wang, Improved data structures for the orthogonal range successor problem, Comput. Geom. 44 (3) (2011) 148–159.